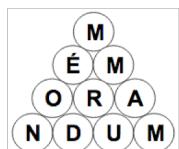


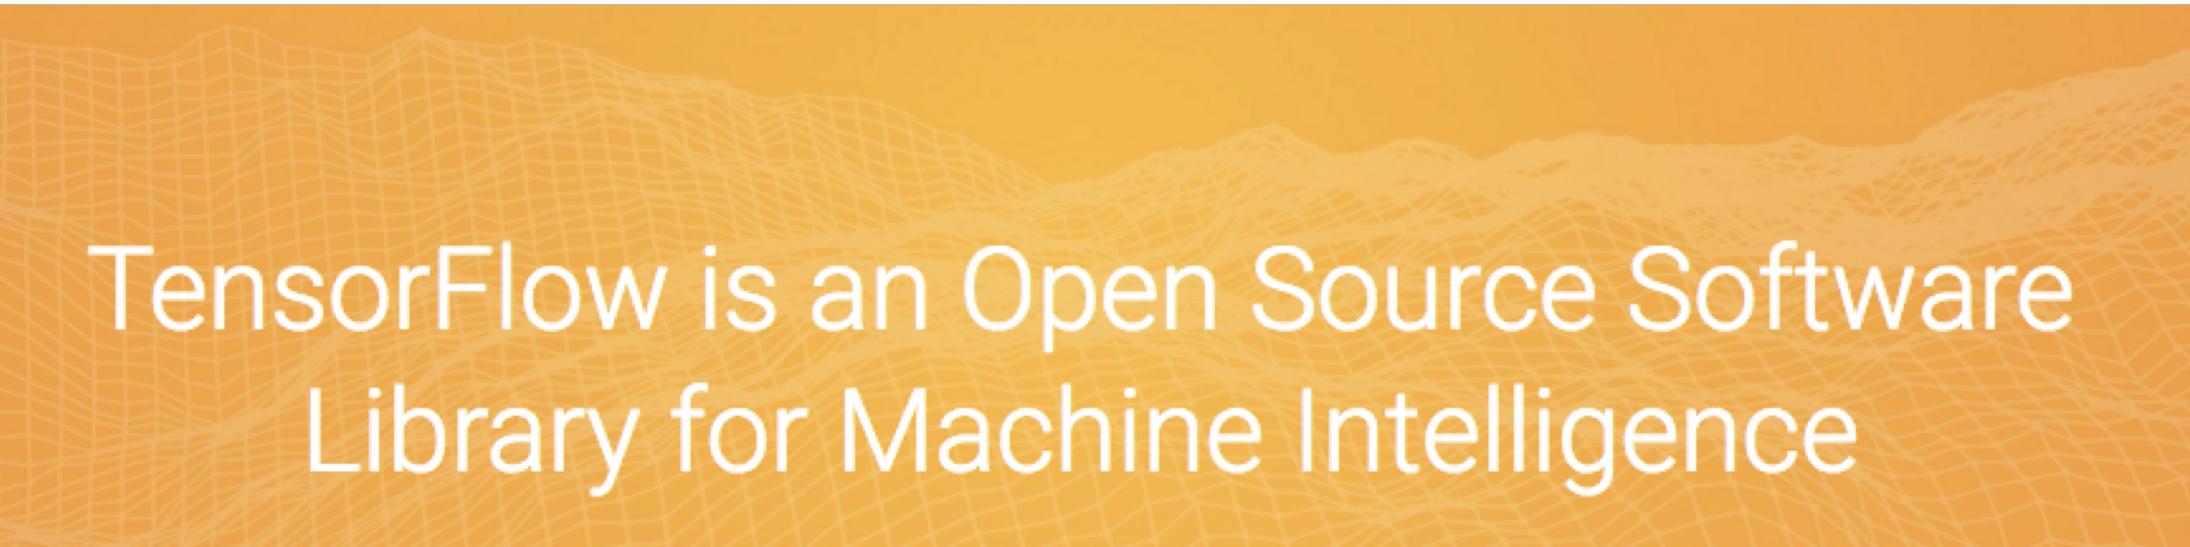


meetup®

Tensorflow
Présentation du MOOC de Google
et Première vue sur l'API

Romain Jouin





The background of the slide features a warm orange gradient. Overlaid on this gradient is a light gray wireframe mesh that forms a series of undulating hills or waves, suggesting a complex data structure or landscape.

TensorFlow is an Open Source Software
Library for Machine Intelligence



Deep Learning

Take machine learning to the next level



Advanced



Approx. 3 months

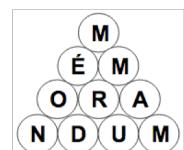
Assumes 6hrs/wk (work
at your own pace)



Join 61,607 students

Built by

Google



▼ Prerequisites and Requirements

This is an intermediate to advanced level course. Prior to taking this course, and in addition to the prerequisites and requirements outlined for the Machine Learning Engineer Nanodegree program, you should possess the following experience and skills:

- Minimum 2 years of programming experience (preferably in Python)
- Git and GitHub experience (assignment code is in a GitHub repo)
- Basic machine learning knowledge (especially supervised learning)
- Basic statistics knowledge (mean, variance, standard deviation, etc.)
- Linear algebra (vectors, matrices, etc.)
- Calculus (differentiation, integration, partial derivatives, etc.)

scipy, numpy, scikit learn...

▼ Syllabus

Lesson 1: From Machine Learning to Deep Learning

- Understand the historical context and motivation for Deep Learning.
- Set up a basic supervised classification task and train a black box classifier on it.
- Train a logistic classifier "by hand" Optimize a logistic classifier using gradient descent, SGD, Momentum and AdaGrad.

Lesson 2: Deep Neural Networks

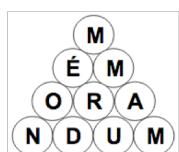
- Train a simple deep network.
- Effectively regularize a simple deep network.
- Train a competitive deep network via model exploration and hyperparameter tuning.

Lesson 3: Convolutional Neural Networks

- Train a simple convolutional neural net.
- Explore the design space for convolutional nets.

Lesson 4: Deep Models for Text and Sequences

- Train a text embedding model.
- Train a LSTM model.



Environnement pour la soirée

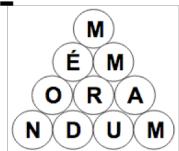


THE MNIST DATABASE of handwritten digits

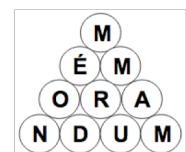
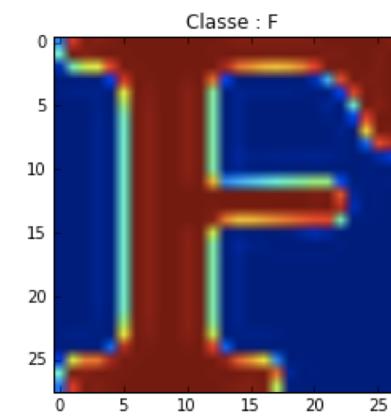
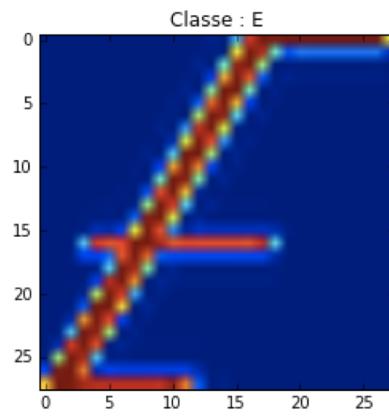
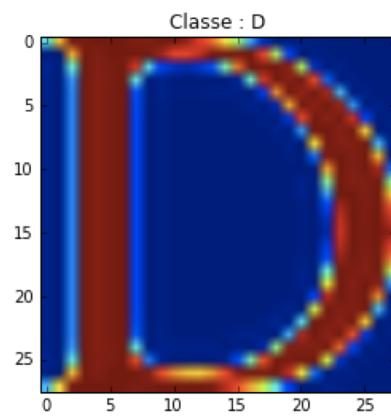
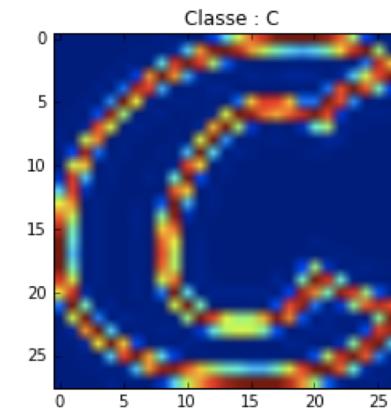
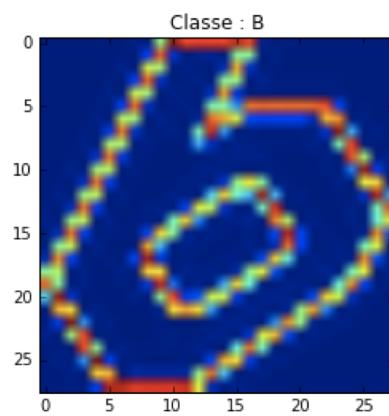
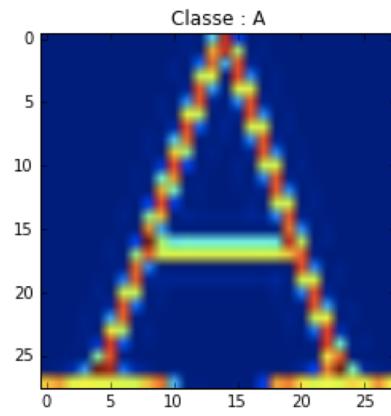
[Yann LeCun](#), Courant Institute, NYU
[Corinna Cortes](#), Google Labs, New York
[Christopher J.C. Burges](#), Microsoft Research, Redmond

The MNIST database of handwritten digits, available from this page, has a training set of 60,000 examples, and a test set of 10,000 examples. It is a subset of a larger set available from NIST. The digits have been size-normalized and centered in a fixed-size image.

<http://yann.lecun.com/exdb/mnist/>



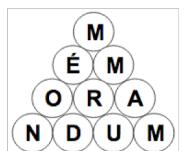
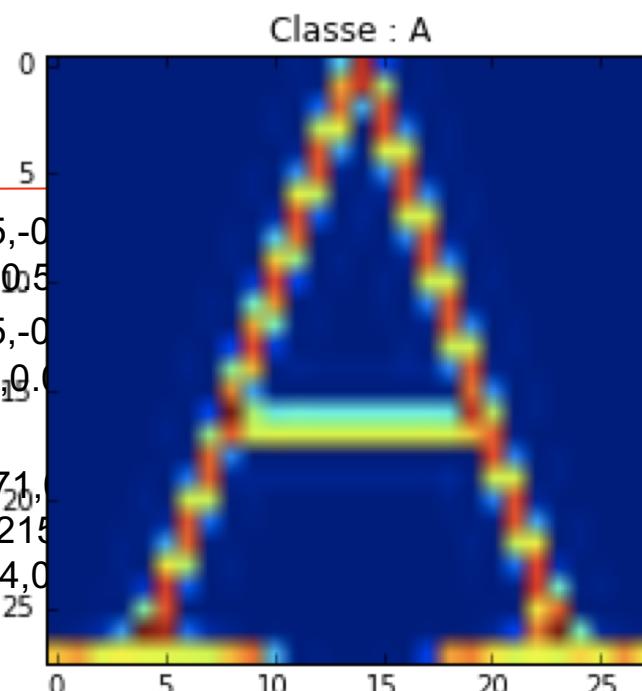
MNIST Database



```
dataset.shape = (52909, 28, 28)
```

```
dataset[0].shape = (28, 28)
```

array(



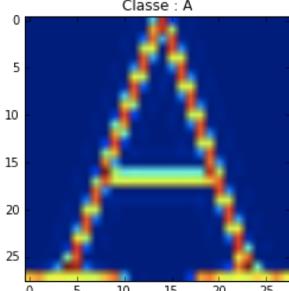
$$\begin{aligned} & \frac{\partial}{\partial w_1} \left[\frac{1}{2} \left(y - \hat{y} \right)^2 + \left(w_1^T x + b \right)^2 \right] \\ &= \frac{\partial}{\partial w_1} \left[\frac{1}{2} \left(y - \frac{e^{w_1^T x + b}}{1 + e^{w_1^T x + b}} \right)^2 + \left(w_1^T x + b \right)^2 \right] \end{aligned}$$



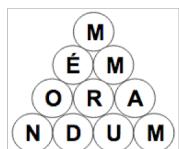
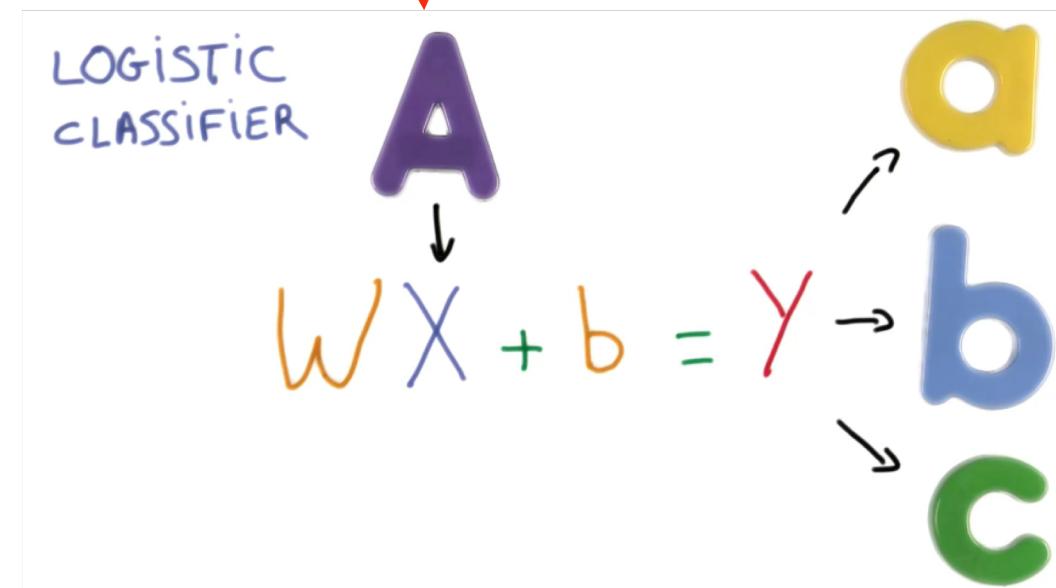
12 millions de paramètres



Logistic classifier et softmax



$$28 * 28 = 756$$



$$\frac{3a(y+z)^2}{39} + \dots$$

$$\frac{a^2(z-y+A)^3}{39} + \dots$$

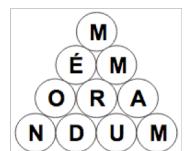
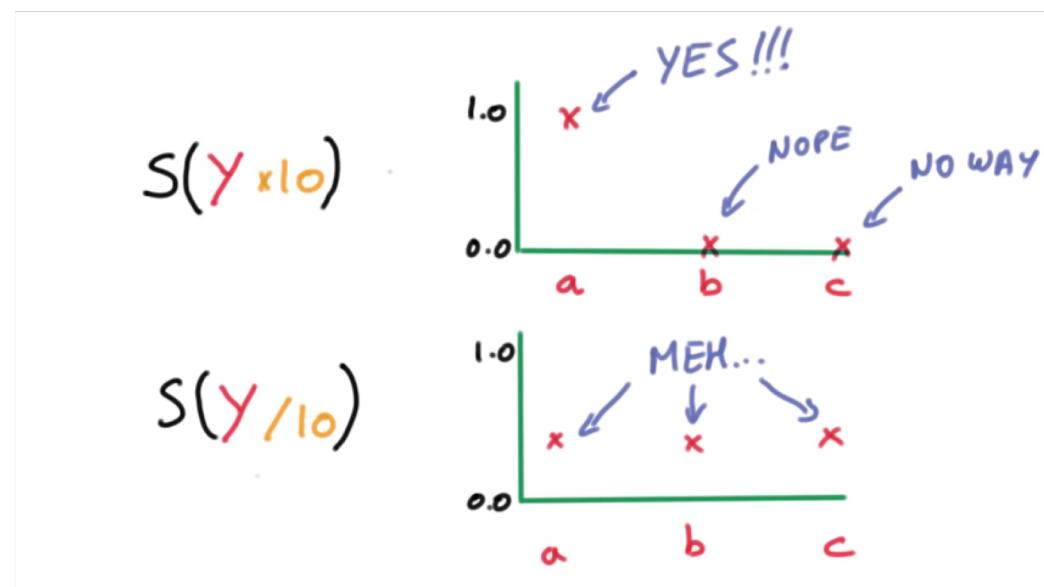
One hot encoding

Pour prendre une décision binaire sur la classification, on peut transformer les probabilités en valeurs "0" ou "1". C'est le "One hot encoding".

Il faut alors un vecteur par classe, qui contiendra plein de zéros, et une seule valeur "1".

Ainsi chaque classe a son vecteur signature qui lui permet d'être reconnaissable.

Le problème du one-hot encoding c'est que ça créé des matrices très "creuses" s'il y a beaucoup de classes : on a quasiment que des zéros, pour un seul "1" par colonne.



$$\begin{aligned}
 & 3a(y+1)^2 + (3y+4)(x+1) \\
 & \frac{a^2(3)}{39} (y+1)^3 + \frac{2}{3}(x+1)^2
 \end{aligned}$$

Cross entropy

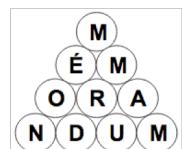
Cette fonction permet de transformer l'ensemble des probabilités en un vecteur binaire avec un seul "1", qui sera la classe éluée.

On a donc :

input => modèle => score => softmax
=>probabilité => cross-entropy => classe choisie.

C'est la classification logistique multinomiale.

$a \rightarrow$	0	0	0	0	0	0
$b \rightarrow$	0	0	0	0	0	0
$c \rightarrow$	0	0	0	0	0	0
$d \rightarrow$	0	0	0	0	0	0
\vdots	0	0	0	0	0	0
.	0	0	0	0	0	0
0	0	0	0	0	0	0

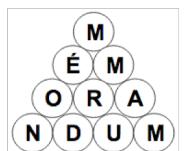
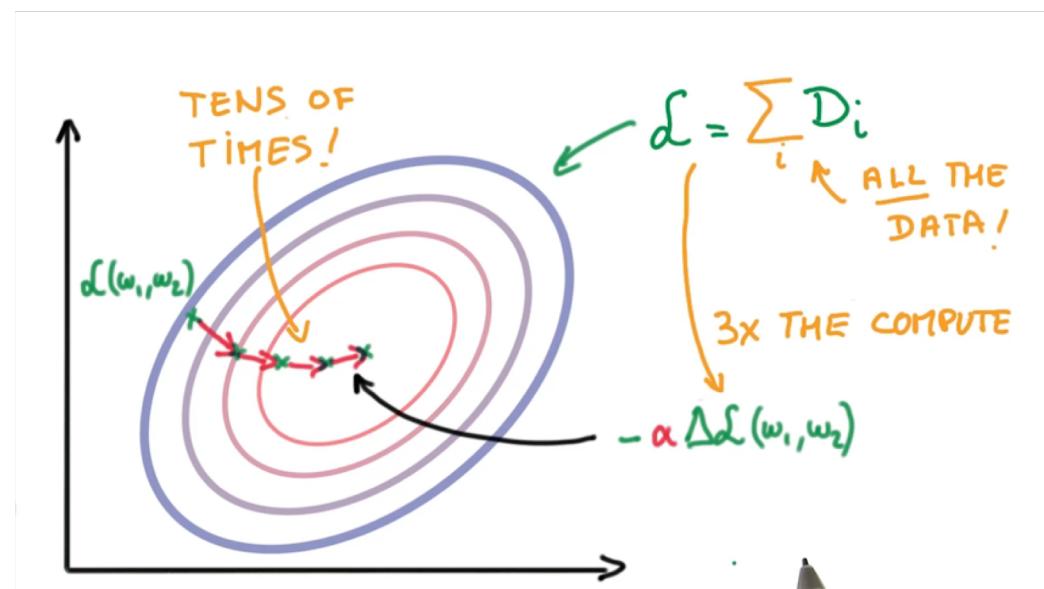


$$\frac{\partial}{\partial w} \left[\frac{1}{2} (y + A)^2 + \frac{1}{3} (y + A)^3 \right]$$

Schotastic gradient descent (vidéo 36)

Comment trouver la matrice des poids [W] et la matrice de biais [b] qui nous permettent d'atteindre notre objectif, à savoir :

- 1) d'avoir une petite distance aux classes correctes
- 2) d'avoir une grande distances aux classes incorrectes



Jeu de donnée utilisé pendant le MOOC



Téléchargement
10 classes : A -> J



training set : 500k images
train_datasets

notMNIST_large/A.pickle
notMNIST_large/B.pickle
notMNIST_large/C.pickle
notMNIST_large/D.pickle
notMNIST_large/E.pickle
notMNIST_large/F.pickle
notMNIST_large/G.pickle
notMNIST_large/H.pickle
notMNIST_large/I.pickle
notMNIST_large/J.pickle

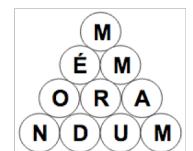
Test set : 19k images
test_datasets

notMNIST_small/A.pickle
notMNIST_small/B.pickle
notMNIST_small/C.pickle
notMNIST_small/D.pickle
notMNIST_small/E.pickle
notMNIST_small/F.pickle
notMNIST_small/G.pickle
notMNIST_small/H.pickle
notMNIST_small/I.pickle
notMNIST_small/J.pickle

Chaque répertoire
contient environ
50 000 images



```
[ 'notMNIST_small/A.pickle',
  'notMNIST_small/B.pickle',
  'notMNIST_small/C.pickle',
  'notMNIST_small/D.pickle',
  'notMNIST_small/E.pickle',
  'notMNIST_small/F.pickle',
  'notMNIST_small/G.pickle',
  'notMNIST_small/H.pickle',
  'notMNIST_small/I.pickle',
  'notMNIST_small/J.pickle']
```

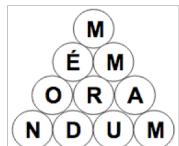


Lesson 1 – classification logistique et descente de gradient

Lesson 2 - Réseau de neurone profond

Lesson 3 - Réseau de convolution

Lesson 4 – Réseau de neurone profond pour l'analyse de texte et de séquences (to do)

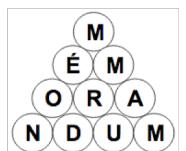
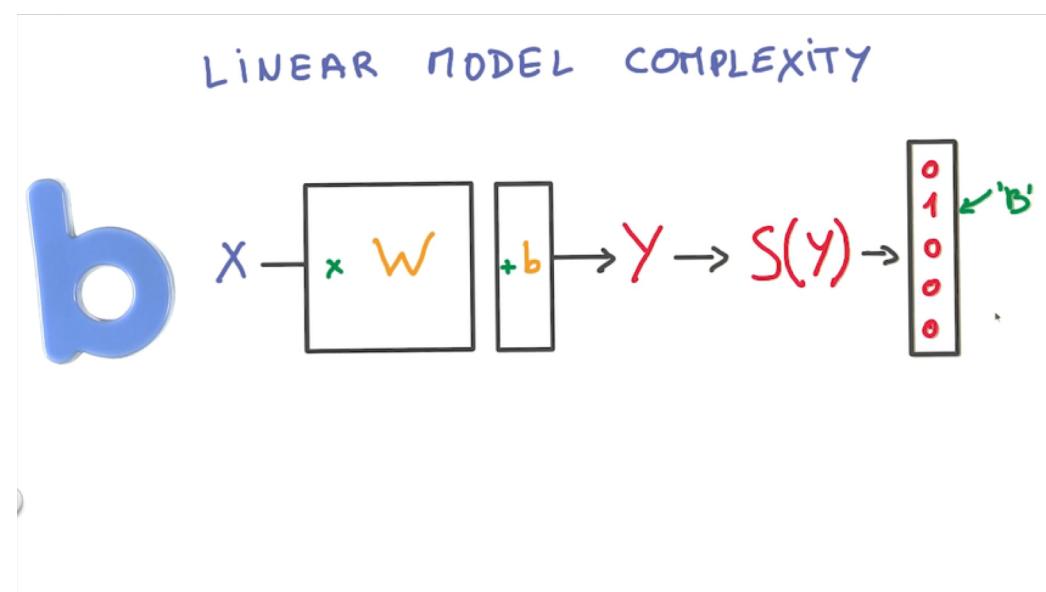


$$3a(y+z)^2 + (3y + 4z + 4A(x+z))^2$$

$$\frac{a^2(3z^2 - 1)}{39} + (y+A)^2 + 2(x+A)^2$$

Nombre de paramètres (vidéo 2)

Les fonctions linéaires sont très stables, et se calculent bien sur GPU. On voudrait donc les conserver. Par contre elles ne permettent pas de faire tout type de calculs. Donc on va chercher à les modifier pour étendre le champ des possibles.



RELU (vidéo 5)

Le « relu » est une fonction très simple : c'est le maximum entre 0 et x

Rappel - dérivée :

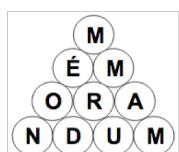
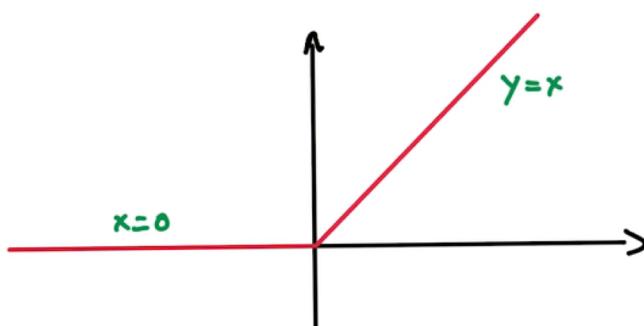
Définition formelle [\[modifier \]](#) [\[modifier le code \]](#)

Soit f une [fonction réelle](#) à valeurs réelles définie sur une réunion quelconque d'[intervalles](#) non triviaux, et x_0 appartenant à l'intérieur de l'ensemble de définition \mathcal{D}_f .

Pour tout $h \in \mathbb{R}^*$ tel que $[x_0, x_0 + h] \subset \mathcal{D}_f$, on appelle *taux d'accroissement de f en x_0* et avec un pas de h la quantité :

$$t_{x_0}(h) = \frac{f(x_0 + h) - f(x_0)}{h}$$

RECTIFIED LINEAR UNITS (RELU)



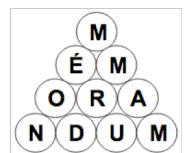
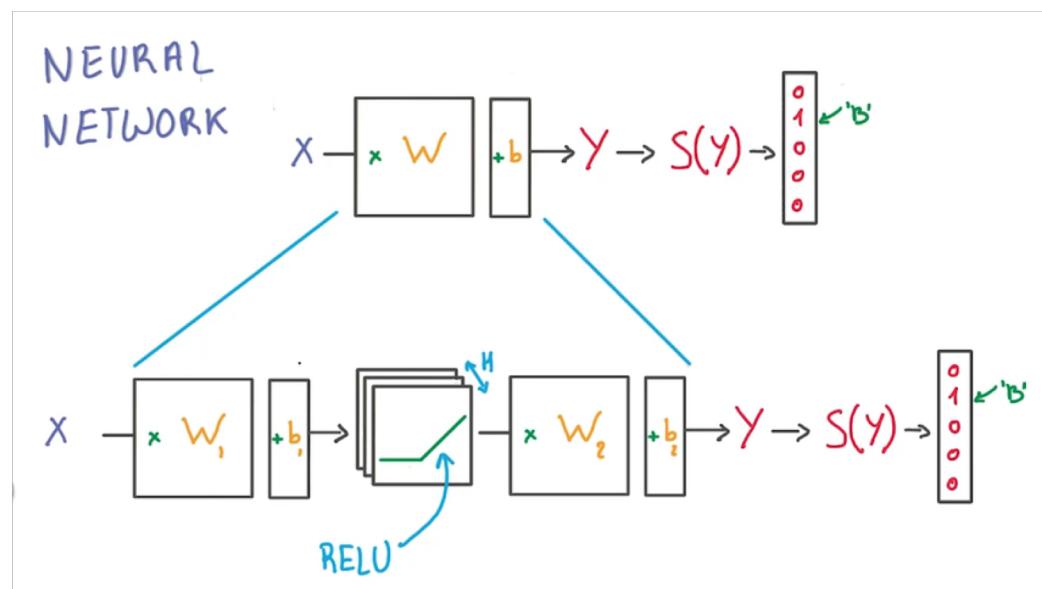
Réseau de neurone profond (vidéo 7)

On rend les fonctions non linéaires en insérant une fonction « relu » au milieu de la fonction...

Les calculs des dérivées restent alors très simple :

$$[g(f(x))]' = g'(f(x)) * f'(x)$$

Comme la descente de gradient fait principalement du calcul de dérivée, il est important que cela reste simple.



$$\frac{\partial}{\partial A} \left[3a(y+z)^2 + (y+A)^3 + 4A(x+z) \right] = 3a(2(y+z)) + 3(y+A)^2 + 4(x+z)$$

$$= 6a(y+z) + 3(y+A)^2 + 4(x+z)$$

$$= 39$$

Back Propagation

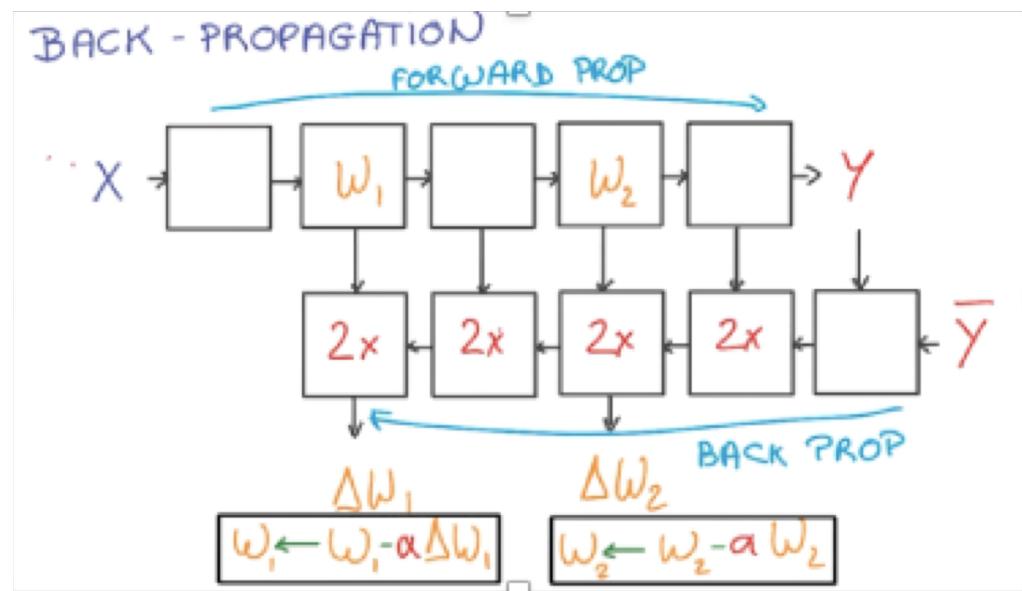
Le calcul du gradient pour un ensemble de poids \mathbf{W} donné se fait en deux étapes :

1) Forward propagation

1) relu -> transformation linéaire -> relu -> transformation linéaire

1) On obtient alors une première prédiction Y

2) Pour affiner cette prédiction et faire la descente de gradient, il faut trouver les dérivées. Pour trouver les dérivées, on va passer par l'étape 2

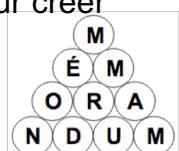


2) Back propagation

Le graph de back-propagation est calculé par Tensorflow (il est déduit du premier graph), on n'a donc pas à le programmer.

Il calcule des dérivées pour de la matrice de poids \mathbf{W} que l'on va soustraire de \mathbf{W} pour créer une nouvelle matrice qui la remplacera.

3) on applique un coefficient alpha sur la dérivée de W avant la soustraction



$$\frac{\partial}{\partial z} [3a(y+z)^2 + 3y + 4A(x+z)]$$

Back Propagation

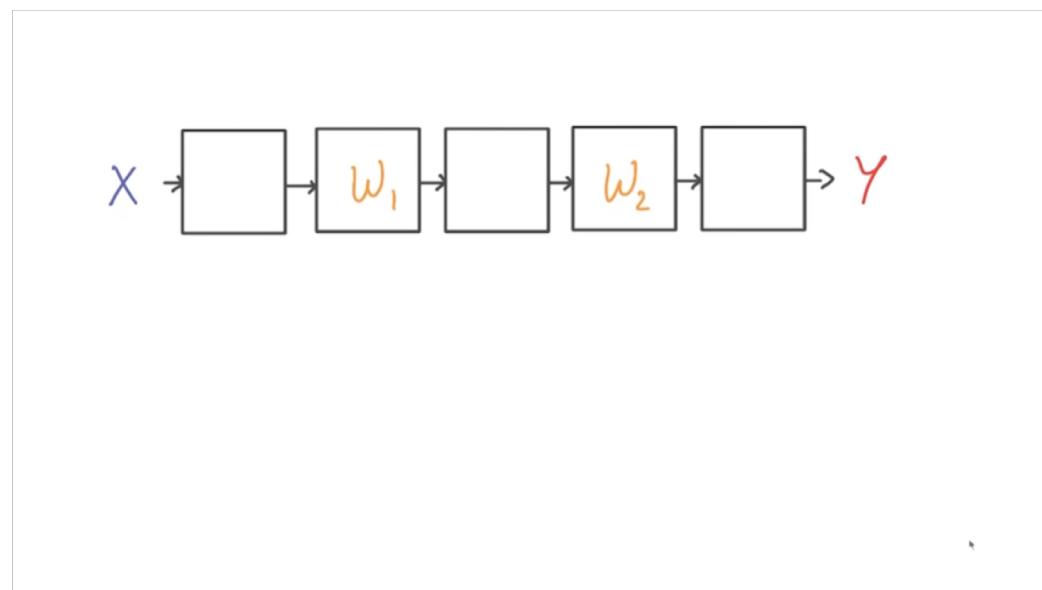
Le calcul du gradient pour un ensemble de poids \mathbf{W} donné se fait en deux étapes :

1) Forward propagation

1) relu -> transformation linéaire -> relu -> transformation linéaire

1) On obtient alors une première prédiction Y

2) Pour affiner cette prédiction et faire la descente de gradient, il faut trouver les dérivées. Pour trouver les dérivées, on va passer par l'étape 2

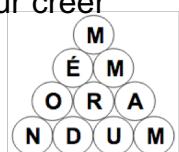


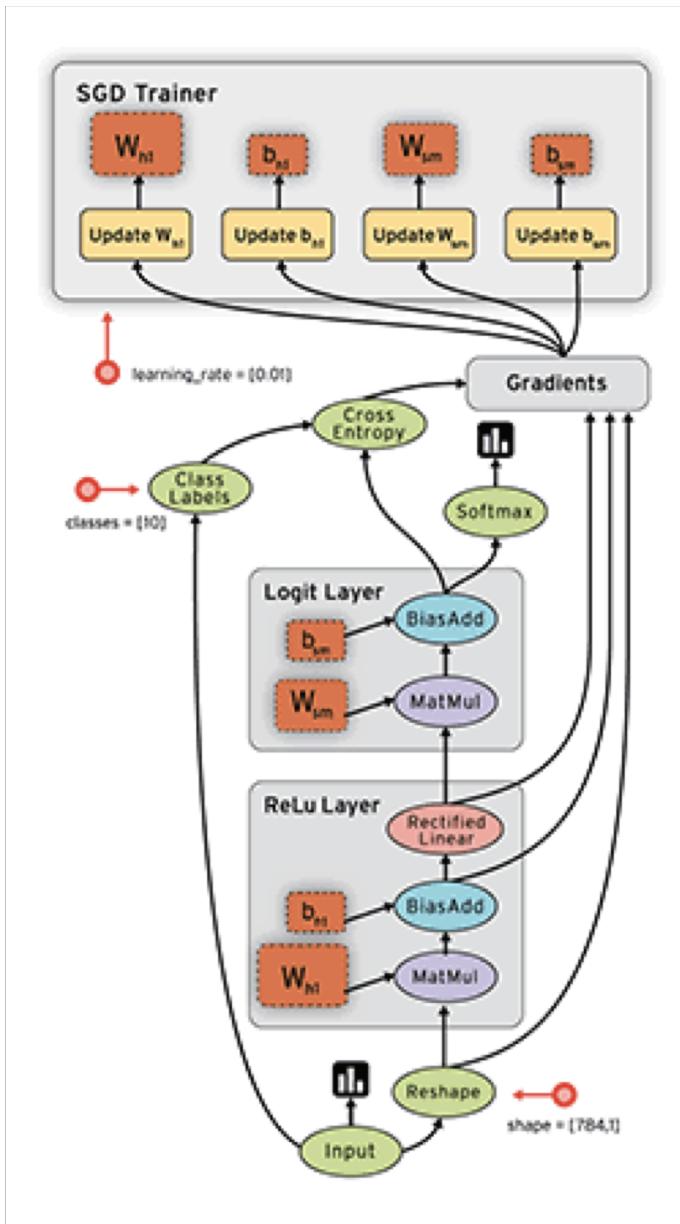
2) Back propagation

Le graph de back-propagation est calculé par Tensorflow (il est déduit du premier graph), on n'a donc pas en le programmer.

Il calcule des dérivées pour de la matrice de poids \mathbf{W} que l'on va soustraire de \mathbf{W} pour créer une nouvelle matrice qui la remplacera.

3) on applique un coefficient alpha sur la dérivée de \mathbf{W} avant la soustraction







https://www.tensorflow.org/versions/r0.7/api_docs/python/state_ops.html#variables

Building Graphs:	26
Constants, Sequences, and Random Values:	14
Variables:	32
Tensor Transformations:	31
Math:	80
Control Flow:	26
Images:	35
Sparse Tensors:	11
Inputs and Readers:	30
Data IO (Python functions):	2
Neural Network:	34
Running Graphs:	20
Training:	35
Wraps python functions:	1
Testing:	6
Layers (contrib):	12
Utilities (contrib):	2

```
#Program that makes up some data in two dimensions, and then fits a line to it.

import tensorflow      as tf
import numpy           as np

# Create 100 phony x, y data points in NumPy, y = x * 0.1 + 0.3
x_data = np.random.rand(100).astype(np.float32)
y_data = x_data * 0.1 + 0.3

# Try to find values for W and b that compute y_data = W * x_data + b
W = tf.Variable(tf.random_uniform([1], -1.0, 1.0))
b = tf.Variable(tf.zeros([1]))
y = W * x_data + b

loss    = tf.reduce_mean(tf.square(y - y_data)) # mean squared errors to minimize
optimizer = tf.train.GradientDescentOptimizer(0.5)
train   = optimizer.minimize(loss)

init    = tf.initialize_all_variables()     # Before starting, initialize the variables
sess    = tf.Session()                      # Launch the graph.
sess.run(init)

for step in xrange(201):                   # Fit the line.
    sess.run(train)
    if step % 20 == 0: print(step, sess.run(W), sess.run(b))

# Learns best fit is W: [0.1], b: [0.3]
```



Building Graphs

```
add_to_collection as_dtype           bytes   control_dependencies  
convert_to_tensor                  convert_to_tensor_or_indexed_slices  
device Dimension                   Dtype    get_collection      get_default_graph  
get_seed             Graph  GraphKeys       import_graph_def  
load_op_library     name_scope     NoGradient     op_scope  
Operation            register_tensor_conversion_function  
RegisterGradient RegisterShape    reset_default_graph    Tensor  
TensorShape
```

Constants, Sequences, and Random Values

```
constant fill          linspace ones      ones_like      random_crop  
random_normal        random_shuffle  random_uniform range  
set_random_seed      truncated_normal zeros      zeros_like
```



Variables

all_variables assert_variables_initialized assign assign_add
assign_sub constant_initializer count_up_to device
get_checkpoint_state get_variable get_variable_scope
IndexedSlices **initialize_all_variables** initialize_variables
latest_checkpoint make_template moving_average_variables
random_normal_initializer random_uniform_initializer Saver scatter_add
scatter_sub scatter_update sparse_mask trainable_variables
truncated_normal_initializer uniform_unit_scaling_initializer
update_checkpoint_state Variable variable_op_scope
variable_scope zeros_initializer

Tensor Transformations

boolean_mask cast concat depth_to_space dynamic_partition
dynamic_stitch expand_dims gather pack pad rank
reshape reverse reverse_sequence shape shape_n size
slice space_to_depth split squeeze string_to_number tile
to_bfloat16 to_double to_float to_int32 to_int64 transpose
unique_with_counts unpack



Math

abs accumulate_n add add_n argmax argmin
batch_cholesky batch_matmul batch_matrix_determinant
batch_matrix_inverse batch_matrix_solve
batch_matrix_solve_ls batch_matrix_triangular_solve
batch_self_adjoint_eig ceil cholesky complex complex_abs
conj cos cross diag div edit_distance erf erfc
exp fft2d floor floordiv ifft2d imag inv
invert_permutation lgamma listdiff log matmul
matrix_determinant matrix_inverse matrix_solve
matrix_solve_ls matrix_triangular_solve maximum minimum
mod mul neg pow real reduce_all reduce_any
reduce_max **reduce_mean** reduce_min reduce_prod
reduce_sum round rsqrt scalar_mul segment_max
segment_mean segment_min segment_prod segment_sum
self_adjoint_eig sign sin sparse_segment_mean
sparse_segment_sqrt_n sparse_segment_sqrt_n_grad
sparse_segment_sum sqrt **square** sub transpose
truediv unique unsorted_segment_sum where



Control Flow

```
add_check_numerics_ops Assert check_numerics cond count_up_to
equal greater greater_equal group identity is_finite is_inf
is_nan less less_equal logical_and logical_not
logical_or logical_xor no_op not_equal Print
select tuple verify_tensor_all_finite where
```

Images

```
adjust_brightness adjust_contrast adjust_hue        adjust_saturation
convert_image_dtype      crop_to_bounding_box    decode_jpeg
decode_png              draw_bounding_boxes     encode_jpeg    encode_png
extract_glimpse         flip_left_right       flip_up_down   grayscale_to_rgb
hsv_to_rgb              pad_to_bounding_box    per_image_whitening
random_brightness        random_contrast       random_flip_left_right
random_flip_up_down      random_hue           random_saturation
resize_area              resize_bicubic        resize_bilinear
resize_image_with_crop_or_pad  resize_images
resize_nearest_neighbor  rgb_to_grayscale    rgb_to_hsv
sample_distorted_bounding_box  saturate_cast    transpose_image
```



Sparse Tensors

shape sparse_concat sparse_fill_empty_rows sparse_reorder
sparse_retain sparse_split sparse_tensor_to_dense
sparse_to_dense sparse_to_indicator SparseTensor
SparseTensorValue

Inputs and Readers

batch batch_join decode_csv decode_json_example
decode_raw FIFOQueue FixedLenFeature
FixedLengthRecordReader FixedLenSequenceFeature
IdentityReader limit_epochs match_filenames_once
matching_files parse_example parse_single_example placeholder
QueueBase RandomShuffleQueue range_input_producer
read_file ReaderBase shuffle_batch shuffle_batch_join size
slice_input_producer string_input_producer TextLineReader
TFRecordReader VarLenFeature WholeFileReader

Data IO (Python functions)

tf_record_iterator TFRecordWriter



Neural Network

```
avg_pool          bias_add          compute_accidental_hits conv2d
conv2d_transpose      depthwise_conv2d      dropout elu
embedding_lookup      fixed_unigram_candidate_sampler in_top_k l2_loss
l2_normalize      learned_unigram_candidate_sampler
local_response_normalization log_uniform_candidate_sampler max_pool
max_pool_with_argmax moments      nce_lossrelu      relu6
sampled_softmax_loss  separable_conv2d      sigmoid
sigmoid_cross_entropy_with_logits softmax softmax_cross_entropy_with_logits
softplus softsign sparse_softmax_cross_entropy_with_logits tanh top_k
uniform_candidate_sampler
```

Running Graphs

```
AbortedError      AlreadyExistsError      CancelledError      DataLossError
DeadlineExceededError FailedPreconditionError get_default_session
InteractiveSession      InternalError      InvalidArgumentError
NotFoundError      OpError      OutOfRangeError PermissionDeniedError
ResourceExhaustedError Session UnauthenticatedError     UnavailableError
UnimplementedError      UnknownError
```



Training

AdagradOptimizer AdamOptimizer add_queue_runner
AggregationMethod clip_by_average_norm clip_by_global_norm
clip_by_norm clip_by_value Coordinator exponential_decay
ExponentialMovingAverage export_meta_graph FtrlOptimizer
generate_checkpoint_state_proto global_norm global_step
GradientDescentOptimizer gradients histogram_summary
image_summary import_meta_graph LooperThread
merge_all_summaries merge_summary MomentumOptimizer
Optimizer QueueRunner RMSPropOptimizer
scalar_summary start_queue_runners stop_gradient
summary_iterator SummaryWriter write_graph zero_fraction

Wraps python functions

py_func

Testing

assert_equal_graph_def compute_gradient
compute_gradient_error get_temp_dir is_built_with_cuda
main



Layers (contrib):

assert_same_float_dtype convolution2d fully_connected
l1_regularizer l2_regularizer summarize_activation
summarize_activations summarize_collection summarize_tensor
summarize_tensors xavier_initializer xavier_initializer_conv2d

Utilities (contrib):

constant_value make_tensor_proto



```
tf.train.Optimizer.minimize(loss, global_step=None, var_list=None,  
gate_gradients=1, aggregation_method=None,  
colocate_gradients_with_ops=False, name=None)
```

Add operations to minimize `loss` by updating `var_list`.

This method simply combines calls `compute_gradients()` and `apply_gradients()`. If you want to process the gradient before applying them call `compute_gradients()` and `apply_gradients()` explicitly instead of using this function.

Args:

- `loss`: A `Tensor` containing the value to minimize.
- `global_step`: Optional `Variable` to increment by one after the variables have been updated.
- `var_list`: Optional list of `Variable` objects to update to minimize `loss`. Defaults to the list of variables collected in the graph under the key `GraphKeys.TRAINABLE_VARIABLES`.
- `gate_gradients`: How to gate the computation of gradients. Can be `GATE_NONE`, `GATE_OP`, or `GATE_GRAPH`.
- `aggregation_method`: Specifies the method used to combine gradient terms. Valid values are defined in the class `AggregationMethod`.
- `colocate_gradients_with_ops`: If `True`, try colocating gradients with the corresponding op.
- `name`: Optional `name` for the returned operation.

Returns:

An Operation that updates the variables in `var_list`. If `global_step` was not `None`, that operation also increments `global_step`.