

Dojo Tensorflow

Jiqiong QIU
Romain Jouin
Denis Oblin

Outline

1. Overview of 1st meetup dojo

- a. What we have learned
- b. Assignment 1

2. Goal

- a. Multilayer perceptron
- b. Regularization
- c. Convolutional neural network

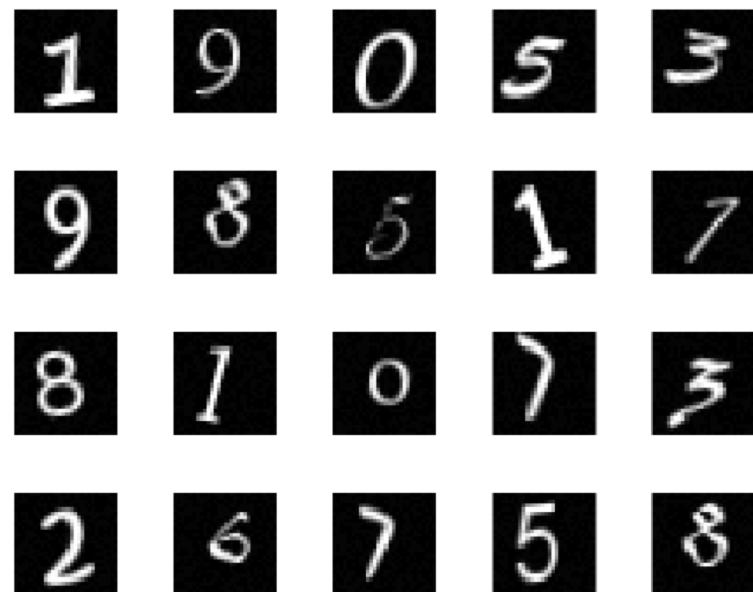
Overview 1st meetup

1. What we have learned

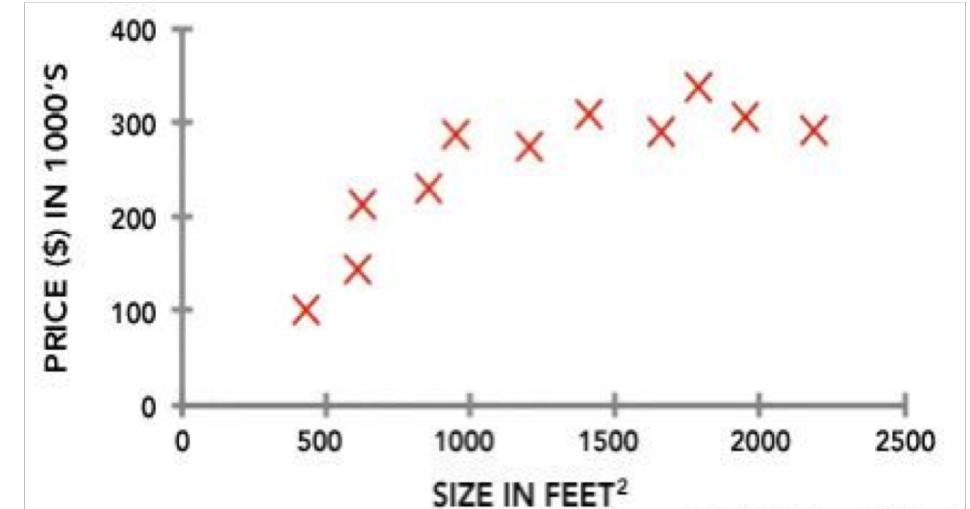
- a. Classification VS Regression
- b. Softmax
- c. One-hot encoding
- d. Cost Function: Cross-entropy

Classification VS Regression

MNIST



Housing Price Prediction



[Credit: Andrew Ng/Stanford]

Softmax function

$$P(y = j | \mathbf{x}) = \frac{e^{\mathbf{x}^\top \mathbf{w}_j}}{\sum_{k=1}^K e^{\mathbf{x}^\top \mathbf{w}_k}}$$

```
f = np.array([123, 456, 789]) # example with 3 classes and each having large scores
p = np.exp(f) / np.sum(np.exp(f)) # Bad: Numeric problem, potential blowup

# instead: first shift the values of f so that the highest number is 0:
f -= np.max(f) # f becomes [-666, -333, 0]
p = np.exp(f) / np.sum(np.exp(f)) # safe to do, gives the correct answer
```

One-hot encoding

Converts a categorical feature to numerical, e.g.,

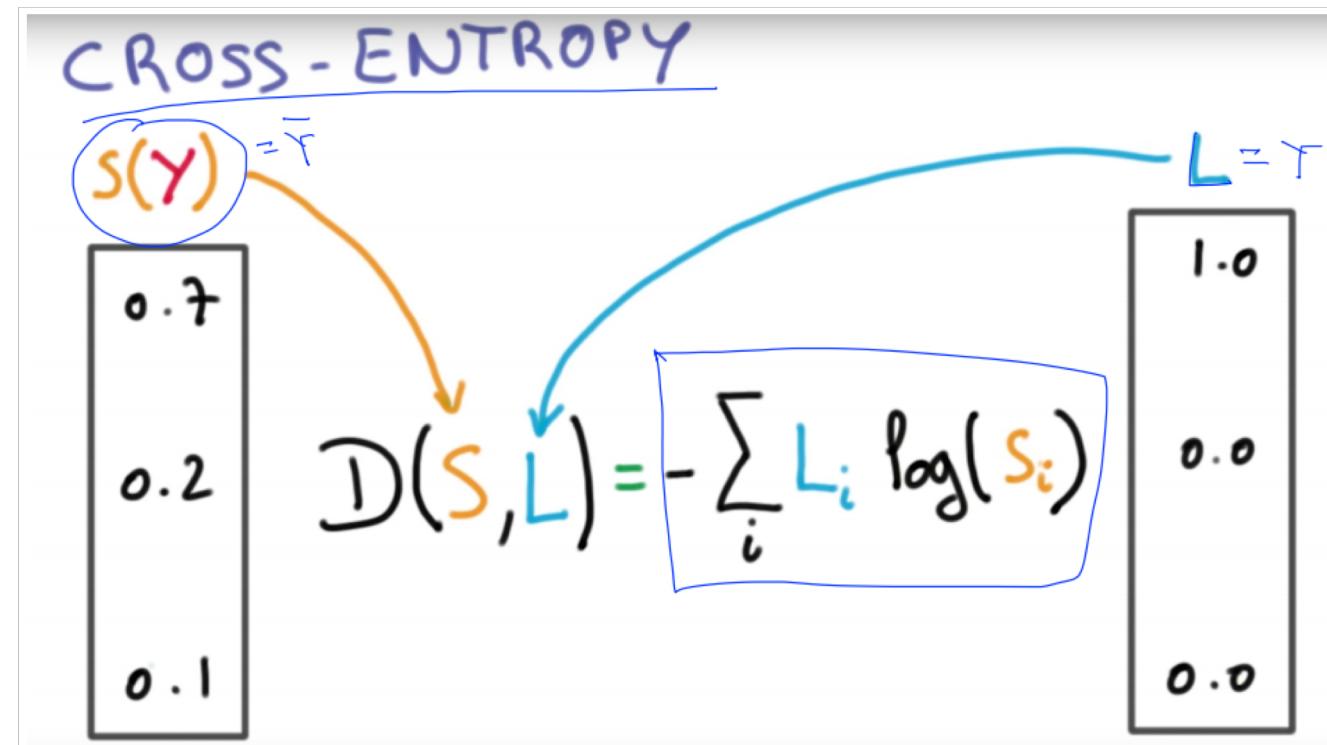
ID	Fruit	apple	banana	watermelon	pear
1	apple	1	0	0	0
2	banana	0	1	0	0
3	banana	0	1	0	0
4	watermelon	0	0	1	0
5	apple	1	0	0	0
6	pear	0	0	0	1
7	pear	0	0	0	1

Cost function: Cross-entropy

Cost function is used to evaluate the performance of learning algorithm:

Sum of squared for regression

Cross-entropy for classification



Outline

1. Overview of 1st meetup dojo

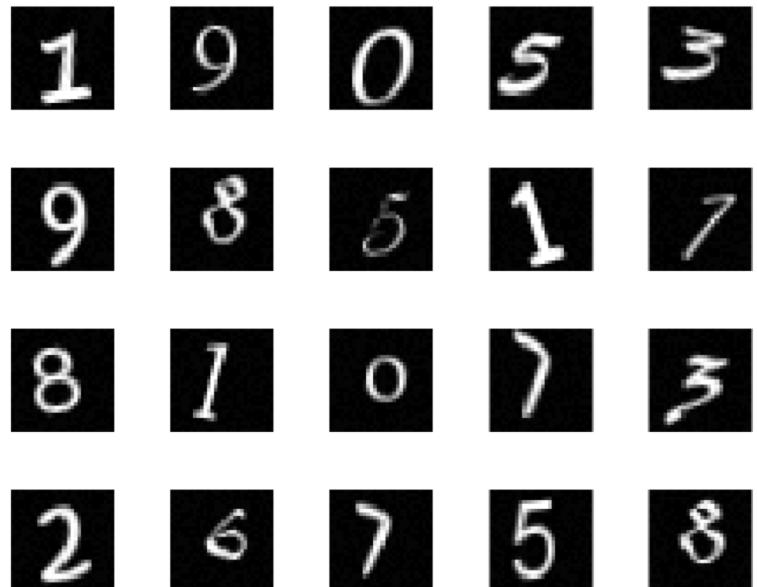
- a. What we have learned
- b. Assignment 1

2. Goal

- a. Multilayer perceptron
- b. Regularization
- c. Convolutional neural network

Assignment 1

NotMNIST dataset : 28x28 pixels, 10 class A - J



Assignment 1

What we have done:

Download the dataset

Display an image/ Study the distribution

Find the duplicated image by using hashing function (MD5)

Realize your first logistic regression model using sklearn

```
print('Training:', train_dataset.shape, train_labels.shape)
print('Validation:', valid_dataset.shape, valid_labels.shape)
print('Testing:', test_dataset.shape, test_labels.shape)
```

What we have now:

3 data sets

```
Training: (200000, 28, 28) (200000,)
Validation: (10000, 28, 28) (10000,)
Testing: (10000, 28, 28) (10000,)
```

Training set: 200 000

Validation set: 10 000

Outline

1. Overview of 1st meetup dojo

- a. What we have learned
- b. Assignment 1

2. Goal

- a. Multilayer perceptron
- b. Regularization
- c. Convolutional neural network

Multilayer perceptron

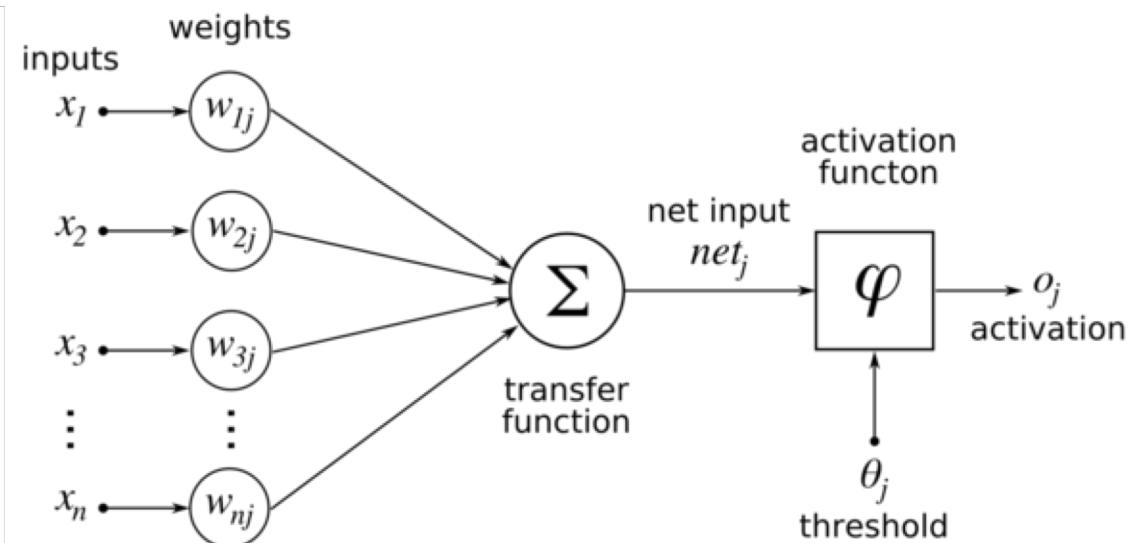
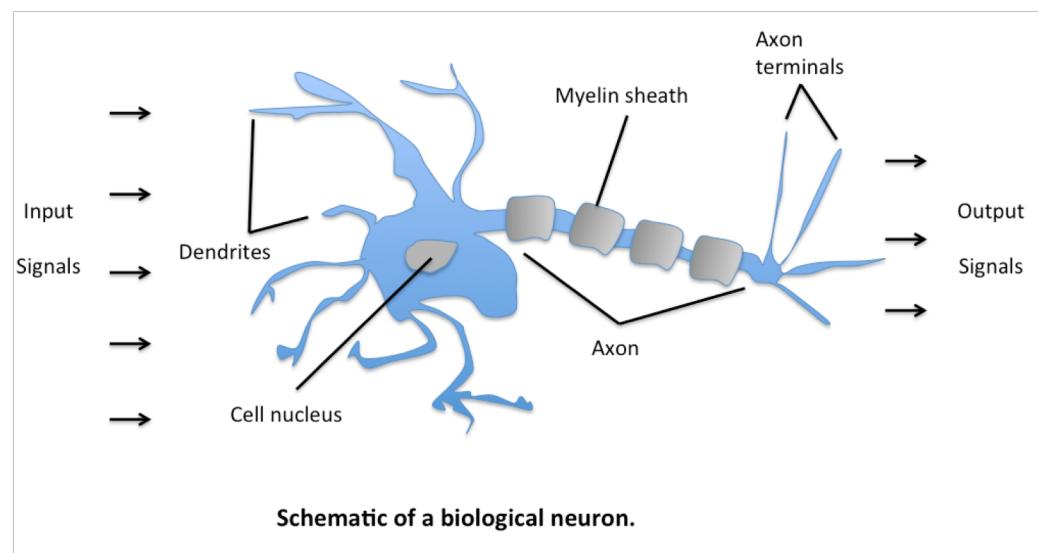
Neuron (perceptron)

Biological neuron and perceptron

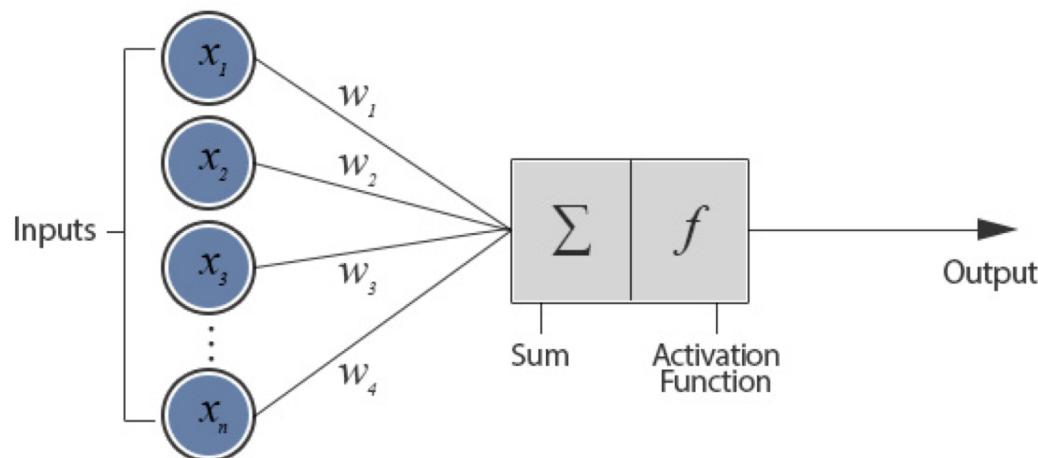
Activation function

MLP (Multilayer perceptron)

Perceptron

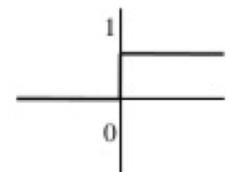


Activation function



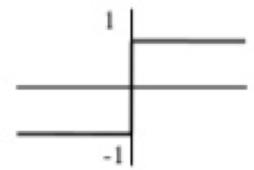
- **Step function**

$$step_t(x) = \begin{cases} 1 & x > t \\ 0 & \text{otherwise} \end{cases}$$



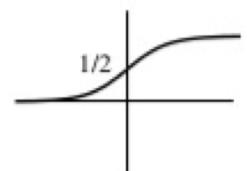
- **Sign function**

$$\text{sign}(x) = \begin{cases} +1 & x \geq 0 \\ -1 & \text{altrimenti} \end{cases}$$

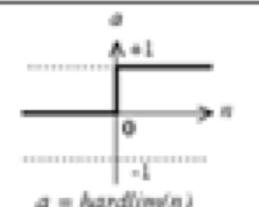
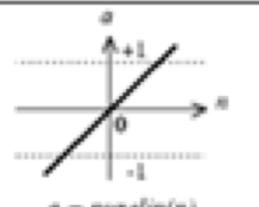
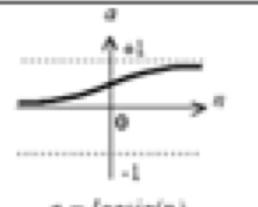
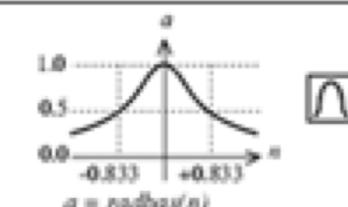
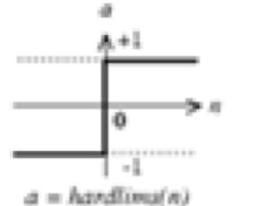
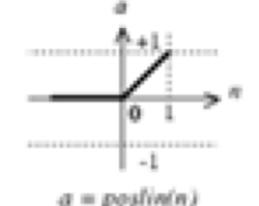
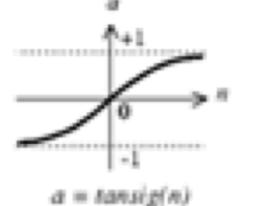
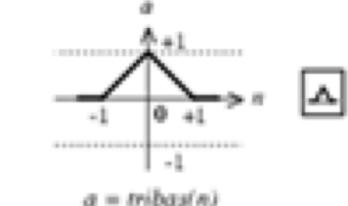
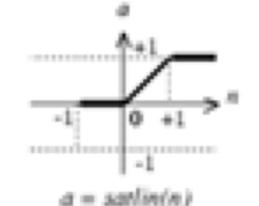
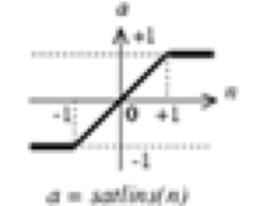
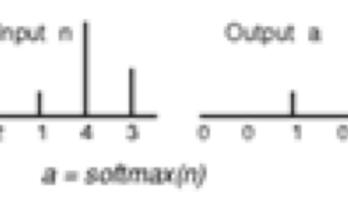
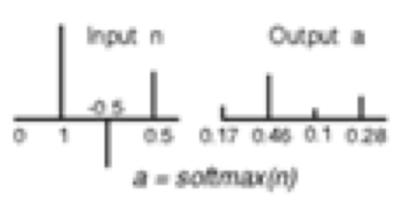


- **Sigmoid function**

$$\text{sigmoide}(x) = \frac{1}{1 + e^{-x}}$$



Activation function

 $a = \text{hardlim}(n)$	 $a = \text{purelin}(n)$	 $a = \text{logsig}(n)$	 $a = \text{radbas}(n)$
Hard-Limit TF  $a = \text{hardlimu}(n)$	Linear TF  $a = \text{poslin}(n)$	Log-Sigmoid TF  $a = \text{tanSig}(n)$	Radial Basis TF  $a = \text{tribas}(n)$
Symmetric Hard-Limit TF  $a = \text{satlin}(n)$	Positive Linear TF  $a = \text{satlins}(n)$	Tan-Sigmoid TF  $a = \text{softmax}(n)$	Triangular Basis TF  $a = \text{softmax}(n)$
Satlin TF	Satlins TF	Compet TF	Softmax TF

Activation function (Tensorflow)

[tf.nn.relu\(features, name=None\)](#)

[tf.nn.relu6\(features, name=None\)](#)

[tf.nn.elu\(features, name=None\)](#)

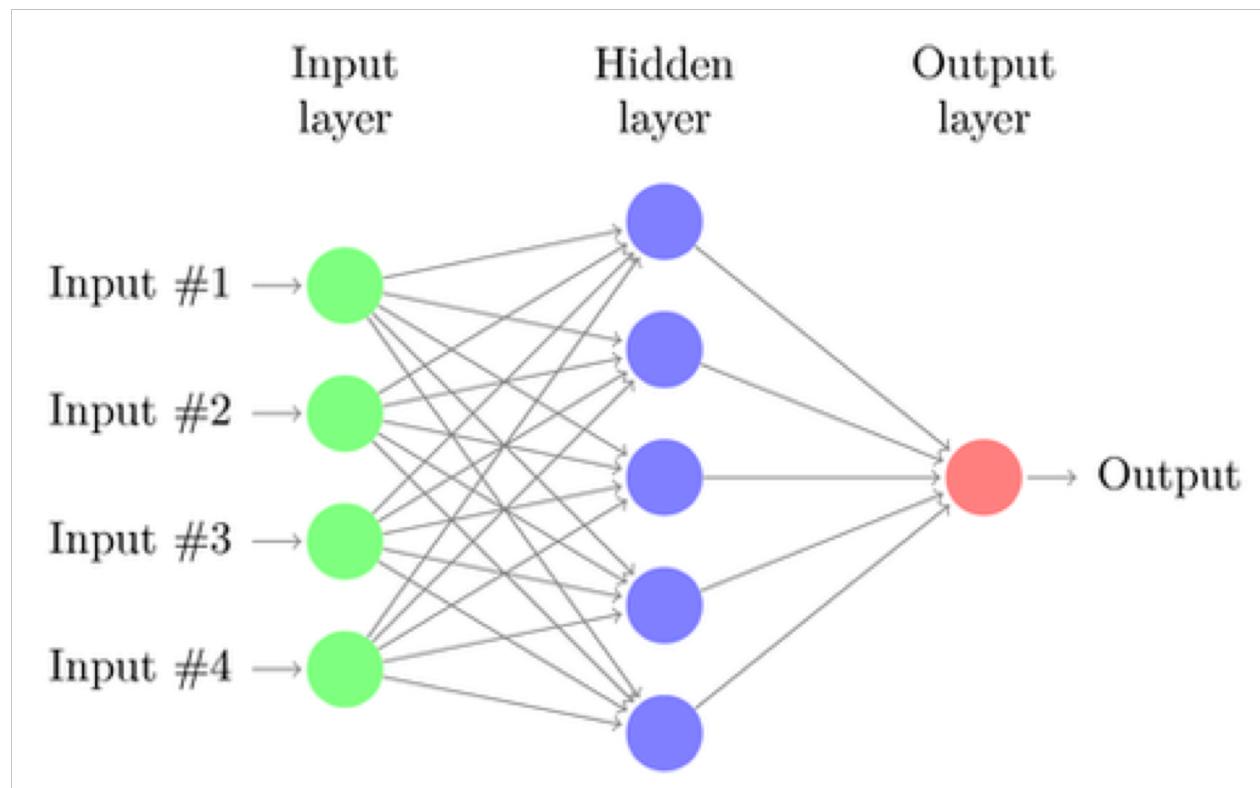
[tf.nn.softplus\(features, name=None\)](#)

[tf.nn.softsign\(features, name=None\)](#)

[tf.nn.dropout\(x, keep_prob, noise_shape=None, seed=None, name=None\)](#)

[tf.nn.bias_add\(value, bias, data_format=None, name=None\)](#)

MLP (Multilayer perceptron)



MLP (Multilayer perceptron)

$$y = \varphi\left(\sum_{i=1}^n w_i x_i + b\right) = \varphi(\mathbf{w}^T \mathbf{x} + b)$$

stretch pixels into single column



input image

0.2	-0.5	0.1	2.0
1.5	1.3	2.1	0.0
0	0.25	0.2	-0.3

W

56
231
24
2

x_i

1.1
3.2
-1.2

b

-96.8
437.9
61.95

$f(x_i; W, b)$

MLP (Multilayer perceptron)

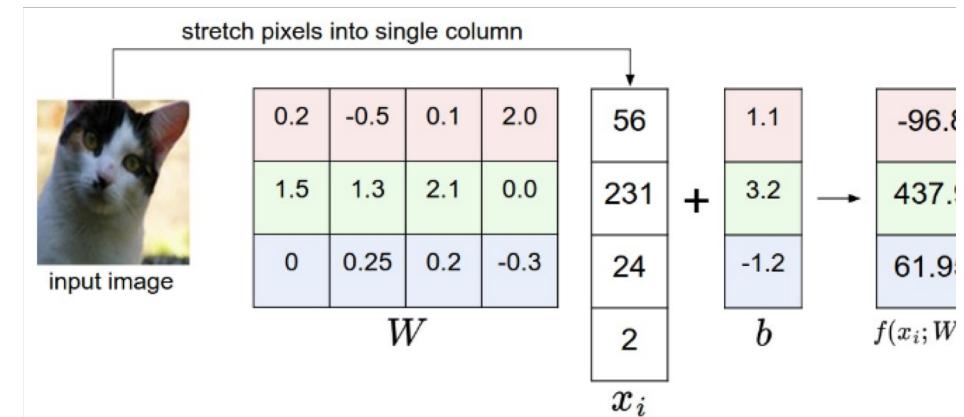
Dataset into constants

```
# Input data.  
# Load the training, validation and test data into constants that are  
# attached to the graph.  
tf_train_dataset = tf.constant(train_dataset[:train_subset, :])  
tf_train_labels = tf.constant(train_labels[:train_subset])  
tf_valid_dataset = tf.constant(valid_dataset)  
tf_test_dataset = tf.constant(test_dataset)
```

Weight and biases defined by variables

```
# Variables.  
# These are the parameters that we are going to be training. The weight  
# matrix will be initialized using random valued following a (truncated)  
# normal distribution. The biases get initialized to zero.  
weights = tf.Variable(  
    tf.truncated_normal([image_size * image_size, num_labels]))  
biases = tf.Variable(tf.zeros([num_labels]))
```

MLP (Multilayer perceptron)



```
# Training computation.  
# We multiply the inputs with the weight matrix, and add biases. We compute  
# the softmax and cross-entropy (it's one operation in TensorFlow, because  
# it's very common, and it can be optimized). We take the average of this  
# cross-entropy across all training examples: that's our loss.  
logits = tf.matmul(tf_train_dataset, weights) + biases  
loss = tf.reduce_mean(  
    tf.nn.softmax_cross_entropy_with_logits(logits, tf_train_labels))
```

MLP (Multilayer perceptron)

```
# Predictions for the training, validation, and test data.  
# These are not part of training, but merely here so that we can report  
# accuracy figures as we train.  
train_prediction = tf.nn.softmax(logits)  
valid_prediction = tf.nn.softmax(  
    tf.matmul(tf_valid_dataset, weights) + biases)  
test_prediction = tf.nn.softmax(tf.matmul(tf_test_dataset, weights) + biases)
```

MLP training

1. Define cost function
2. Find the optimal weight and bias which minimise the cost function

Demo by tensorflow: <http://playground.tensorflow.org>

```
# Optimizer.  
# We are going to find the minimum of this loss using gradient descent.  
optimizer = tf.train.GradientDescentOptimizer(0.5).minimize(loss)
```

MLP training

```
num_steps = 801

def accuracy(predictions, labels):
    return (100.0 * np.sum(np.argmax(predictions, 1) == np.argmax(labels, 1))
            / predictions.shape[0])

with tf.Session(graph=graph) as session:
    # This is a one-time operation which ensures the parameters get initialized as
    # we described in the graph: random weights for the matrix, zeros for the
    # biases.
    tf.initialize_all_variables().run()
    print('Initialized')
    for step in range(num_steps):
        # Run the computations. We tell .run() that we want to run the optimizer,
        # and get the loss value and the training predictions returned as numpy
        # arrays.
        _, l, predictions = session.run([optimizer, loss, train_prediction])
        if (step % 100 == 0):
            print('Loss at step %d: %f' % (step, l))
            print('Training accuracy: %.1f%%' % accuracy(
                predictions, train_labels[:train_subset, :]))
            # Calling .eval() on valid_prediction is basically like calling run(), but
            # just to get that one numpy array. Note that it recomputes all its graph
            # dependencies.
            print('Validation accuracy: %.1f%%' % accuracy(
                valid_prediction.eval(), valid_labels))
    print('Test accuracy: %.1f%%' % accuracy(test_prediction.eval(), test_labels))
```

Assignment 2

Turn the logistic regression example with SGD into a 1-hidden layer neural network with rectified linear units (`nn.relu()`) and 1024 hidden nodes. This model should improve your validation / test accuracy.

Outline

1. Overview of 1st meetup dojo

- a. What we have learned
- b. Assignment 1

2. Goal

- a. Multilayer perceptron
- b. Regularization
- c. Convolutional neural network

Regularization

D

Outline

1. Overview of 1st meetup dojo

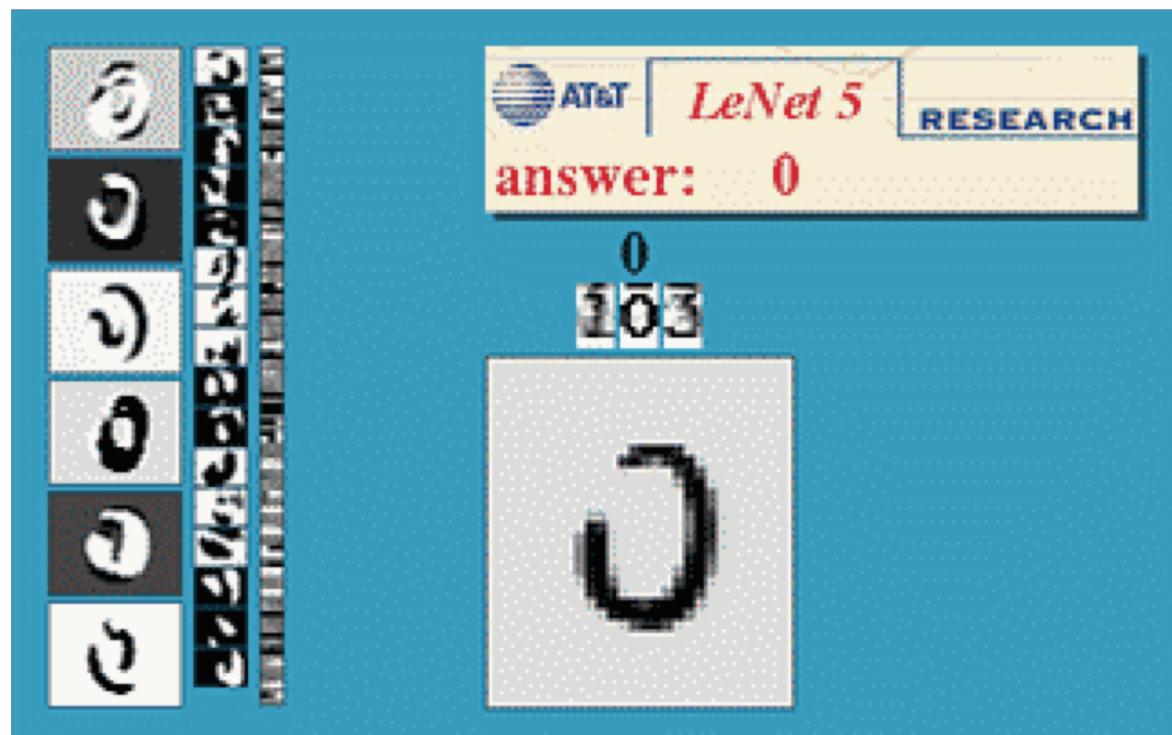
- a. What we have learned
- b. Assignment 1

2. Goal

- a. Multilayer perceptron
- b. Regularization
- c. Convolutional neural network

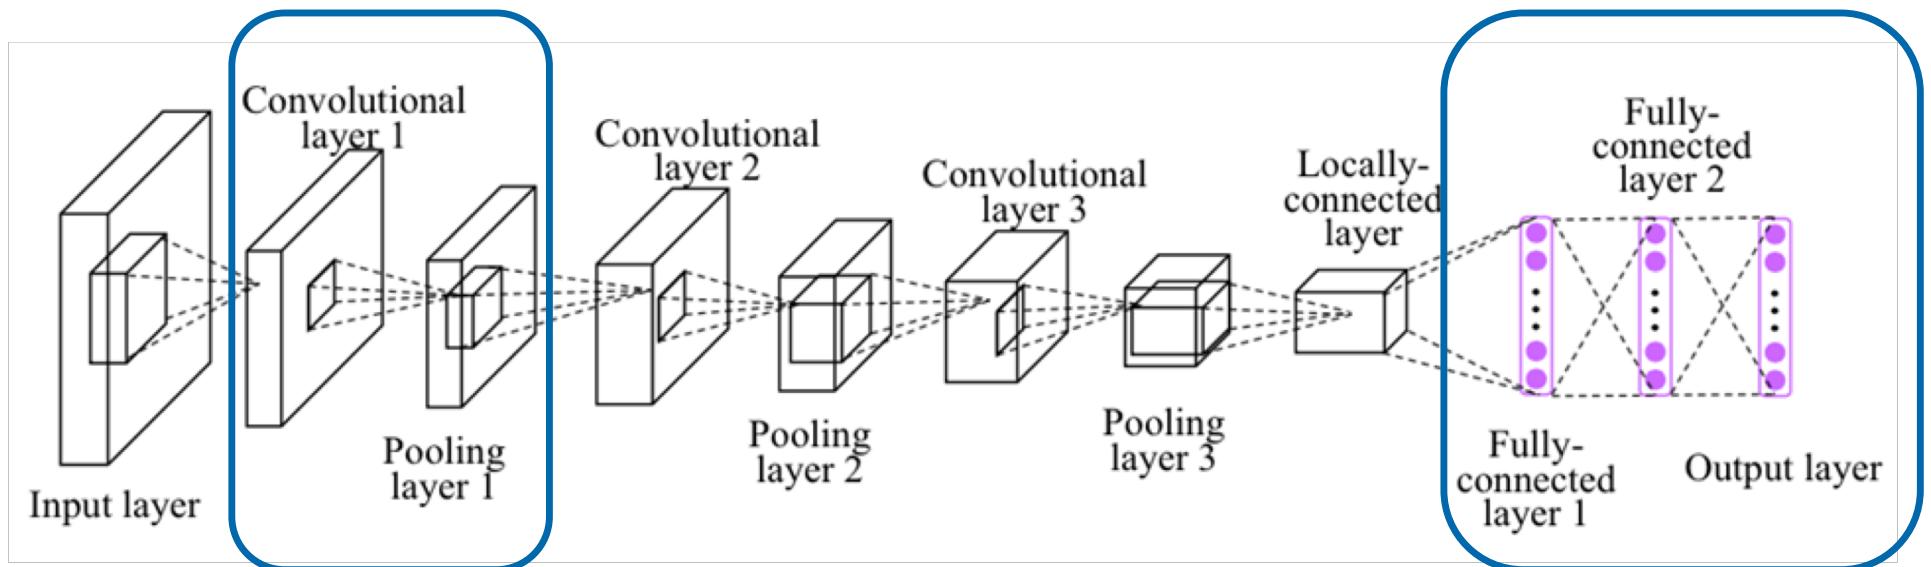
Convolutionnal Neural Network

Yann Lecun



Convolutionnal Neural Network

Convolution + Pooling



Convolutionnal Neural Network

Convolution layer

Pooling layer

MLP and more

Training and Results

Convolution Layer

Natural Images have the property of being “stationary”



Convolution Layer



1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Image

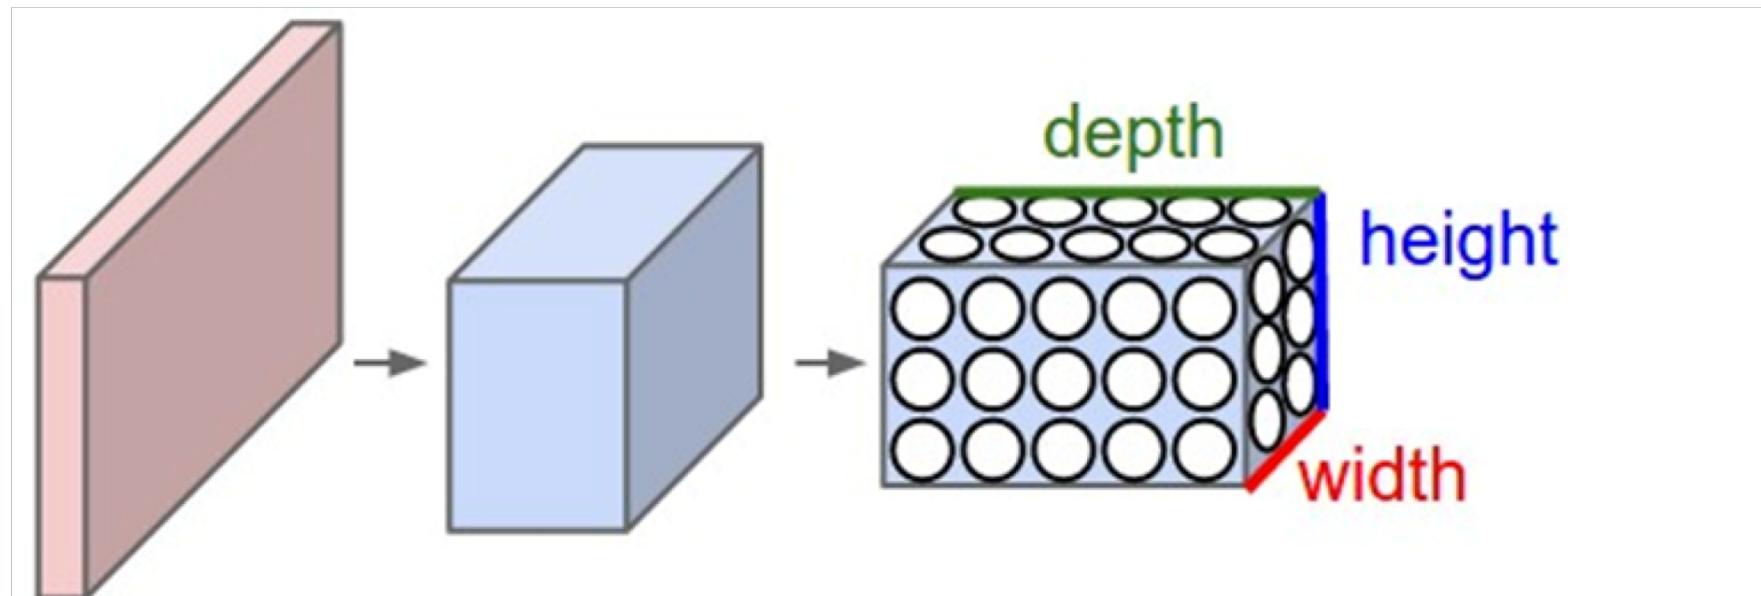
Kernel Matrix

4		

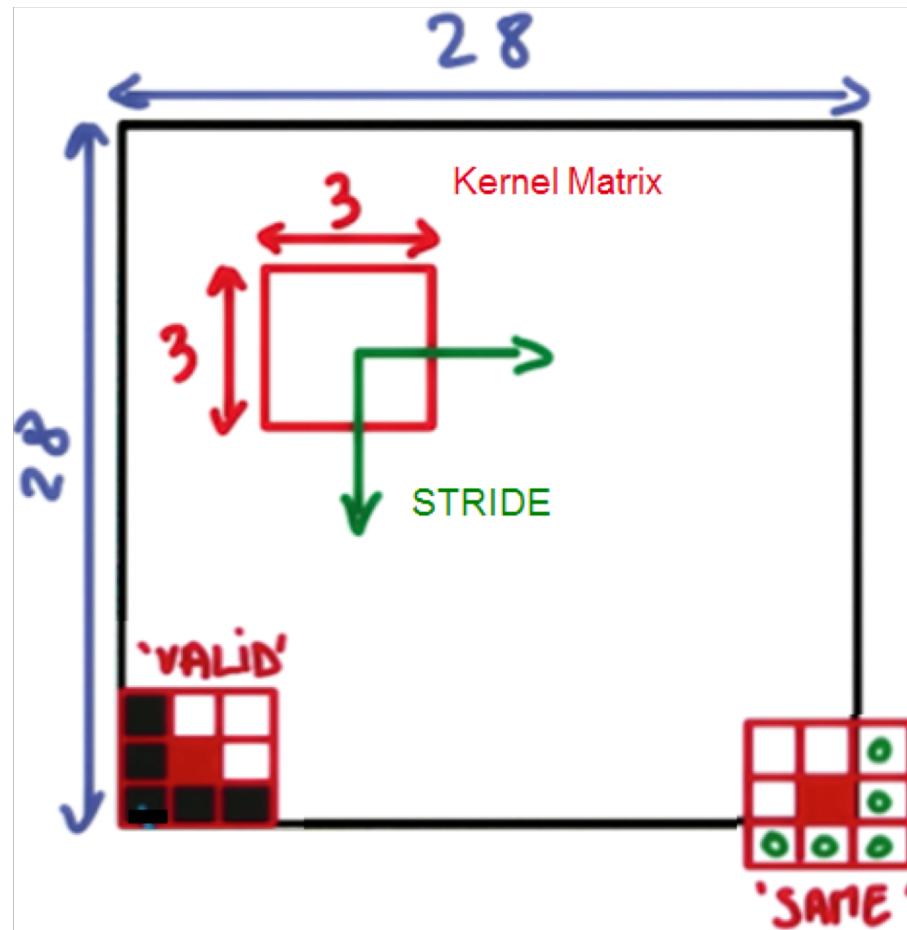
Convolved
Feature

Convolution Layer

Input Depth = 3 (rgb)
Output Depth = 400



Convolution Layer



Convolution Layer

```
def conv2d(x, w):
    return tf.nn.conv2d(x, w, strides=[1, 1, 1, 1], padding='SAME')
```

Usually, a relu function is used

Convolutionnal Neural Network

Convolution layer

Pooling layer

MLP and more

Training and Results

Pooling